

## Experiment 1 : Implementation of Linear and Logistic Regression on Real-World Datasets

**Aim:** Implement Multi Regression, Lasso, and Ridge Regression on real-world datasets.

### Theory:

#### **A) Linear Regression**

##### **1. Dataset Source**

Dataset Name: **Real Estate Valuation Data (Taipei Housing Dataset)**

Source Link:

<https://www.kaggle.com/datasets/hastingssibanda/taipei-housing-dataset-uci>

This dataset is a real-world housing valuation dataset collected in Taipei, Taiwan. It is suitable for linear regression because the target variable is continuous and influenced by multiple numerical predictors.

##### **2. Dataset Description**

The dataset contains records of real estate transactions and is used to predict the **house price per unit area**.

##### **Features include:**

- Transaction date
- House age
- Distance to nearest MRT station
- Number of convenience stores
- Latitude
- Longitude

##### **Target Variable:**

- House price per unit area (continuous value)

##### **Dataset Size:**

- 414 records
- 7 columns (6 input features + 1 output variable)

### Characteristics:

- Numerical dataset
- No complex categorical variables
- Suitable for regression modeling
- Represents real-world property valuation
- Contains geographic and socio-economic factors

This dataset is commonly used for housing price prediction and regression analysis.

### 3. Mathematical Formulation of Linear Regression

Linear Regression models the relationship between independent variables and a continuous dependent variable.

The model is defined as:

$$y = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \dots + \beta_n x_n + \epsilon$$

Where:

- $y$  = predicted output
- $x_i$  = input features
- $\beta_i$  = coefficients
- $\epsilon$  = error term

The model minimizes the **Mean Squared Error (MSE)**:

$$MSE = \frac{1}{n} \sum (y_i - \hat{y}_i)^2$$

The optimal coefficients are estimated using Ordinary Least Squares.

### 4. Algorithm Limitations

Linear Regression has several limitations:

- Assumes a linear relationship between variables
- Sensitive to outliers
- Cannot model complex non-linear patterns
- Affected by multicollinearity
- Requires normally distributed residuals
- Performs poorly with irrelevant or noisy features

It is not suitable for highly non-linear or categorical-heavy datasets without transformation.

## 5. Methodology / Workflow

The experiment follows these steps:

1. Load dataset
2. Data cleaning and preprocessing
3. Feature selection
4. Train-test split
5. Model training using Linear Regression
6. Prediction on test data
7. Model evaluation

### Workflow:

Dataset → Preprocessing → Train/Test Split → Train Model → Predict → Evaluate

## 6. Performance Analysis

The model performance is evaluated using:

- Mean Squared Error (MSE)
- Root Mean Squared Error (RMSE)
- R<sup>2</sup> Score

A low MSE indicates accurate predictions.

RMSE shows average prediction error magnitude.

R<sup>2</sup> score measures how well the model explains variance.

The model demonstrates good predictive capability with acceptable error levels and strong correlation between predicted and actual house prices.

## 7. Hyperparameter Tuning

Linear Regression has limited hyperparameters. Regularization techniques such as Ridge and Lasso can be applied to improve generalization.

Cross-validation was used to verify model stability.

The tuned model reduced overfitting and improved prediction consistency.

## B) Logistic Regression

### 1. Dataset Source

Dataset Name: **Loan Approval Dataset**

Source Link:

<https://www.kaggle.com/datasets/arbaaztamboli/loan-approval-dataset>

This dataset is a real-world financial dataset used to predict whether a loan application is approved or rejected.

### 2. Dataset Description

The dataset contains applicant financial and personal information used to classify loan approval status.

**Features include:**

- Gender
- Marital status
- Income
- Education
- Credit history
- Loan amount
- Employment status

**Target Variable:**

- Loan\_Status (Approved / Rejected)

**Dataset Size:**

- ~600 records
- Multiple categorical and numerical features

**Characteristics:**

- Binary classification dataset
- Contains categorical variables
- Real-world banking scenario
- Some missing values present

### 3. Mathematical Formulation of Logistic Regression

Logistic Regression predicts the probability of class membership using the sigmoid function:

$$P(y = 1) = \frac{1}{1 + e^{-z}}$$

Where:

$$z = \beta_0 + \beta_1 x_1 + \dots + \beta_n x_n$$

The loss function used is Binary Cross Entropy:

$$L = -[y \log(p) + (1 - y) \log(1 - p)]$$

The model is optimized using Gradient Descent.

### 4. Algorithm Limitations

Logistic Regression limitations include:

- Assumes linear decision boundary
- Poor performance on non-linear data
- Sensitive to multicollinearity
- Struggles with highly imbalanced datasets
- Requires proper feature scaling
- Limited performance on complex relationships

### 5. Methodology / Workflow

Steps followed:

1. Load dataset
2. Handle missing values
3. Encode categorical variables
4. Feature scaling
5. Train-test split
6. Logistic Regression training
7. Prediction
8. Evaluation

**Workflow:**

Dataset → Cleaning → Encoding → Scaling → Train Model → Predict → Evaluate

**6. Performance Analysis**

Evaluation metrics used:

- Accuracy
- Precision
- Recall
- F1 Score
- Confusion Matrix

The model achieves high classification accuracy and balanced precision-recall, indicating effective prediction of loan approval status.

The confusion matrix confirms correct separation of approved and rejected applications.

**7. Hyperparameter Tuning**

Hyperparameters tuned:

- Regularization strength (C)
- Solver selection

Grid Search Cross Validation was applied to select optimal values.

Tuning improved model generalization and classification reliability.

**Code:****Linear Regression:**

```
import pandas as pd

import numpy as np

import matplotlib.pyplot as plt

from sklearn.model_selection import train_test_split

from sklearn.linear_model import LinearRegression

from sklearn.metrics import mean_squared_error, r2_score

data = pd.read_excel('/content/Real estate valuation data set.xlsx')

X = data.drop('Y house price of unit area', axis=1)

y = data['Y house price of unit area']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)

model = LinearRegression()

model.fit(X_train, y_train)

y_pred = model.predict(X_test)

print("MSE:", mean_squared_error(y_test, y_pred))

print("R2 Score:", r2_score(y_test, y_pred))
```

**Logistic Regression :**

```
import pandas as pd

from sklearn.model_selection import train_test_split

from sklearn.linear_model import LogisticRegression
```

```
from sklearn.preprocessing import LabelEncoder, StandardScaler # Added
StandardScaler

from sklearn.metrics import accuracy_score, classification_report

loan = pd.read_csv('/content/Loan Dataset.csv')

loan = loan.fillna()

encoder = LabelEncoder()

for column in loan.columns:

    if loan[column].dtype == 'object':

        loan[column] = encoder.fit_transform(loan[column])



X = loan.drop('Loan_Approval_Status', axis=1)

y = loan['Loan_Approval_Status']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)

scaler = StandardScaler()

X_train_scaled = scaler.fit_transform(X_train)

X_test_scaled = scaler.transform(X_test)





model = LogisticRegression(max_iter=200)

model.fit(X_train_scaled, y_train)

y_pred = model.predict(X_test_scaled)

print("Accuracy:", accuracy_score(y_test, y_pred))

print(classification_report(y_test, y_pred))
```


**Output:****Linear Regression:**


```
MSE: 101.29181119765644
R2 Score: 0.538794526802929
```

**Logistic Regression:**

```
Accuracy: 0.8547115384615385
precision    recall   f1-score   support
0            0.84     0.72      0.77     3631
1            0.86     0.93      0.89     6769

accuracy          0.85
macro avg       0.85     0.82      0.83     10400
weighted avg    0.85     0.85      0.85     10400
```





### **Conclusion :**

Linear Regression effectively predicts continuous housing prices, while Logistic Regression successfully classifies loan approval outcomes.

Both algorithms demonstrate strong performance on real-world datasets when proper preprocessing and tuning are applied. The experiments highlight the importance of choosing suitable models based on problem type (regression vs classification).