

Experiment 8: Design and train a Convolutional Neural Network (CNN) on image datasets (e.g., MNIST/CIFAR)

Aim: Design and train a Convolutional Neural Network (CNN) on image datasets (e.g., MNIST/CIFAR).

Theory:

1. Dataset Source

Dataset used: **MNIST Handwritten Digit Dataset**

Source: TensorFlow Datasets

Link: <https://www.tensorflow.org/datasets/catalog/mnist>

2. Dataset Description

The MNIST dataset is a widely used benchmark dataset for image classification tasks.

Features:

- Each image is a **28 × 28 grayscale image**
- Pixel values range from **0 to 255**

Target Variable:

- Digit label (0 to 9)

Dataset Size:

- Training samples: **60,000 images**
- Testing samples: **10,000 images**

Characteristics:

- Simple and clean dataset
- Suitable for beginners in deep learning
- Used for handwritten digit recognition

3. Mathematical Formulation of CNN

Convolution Operation:

$$\text{Feature Map} = \text{Input} * \text{Kernel}$$

Neuron Output:

$$Z = \sum (w_i x_i) + b$$

$$A = f(Z)$$

Activation Function (ReLU):

$$f(x) = \max(0, x)$$

Softmax Function (Output Layer):

$$P(y_i) = \frac{e^{z_i}}{\sum e^{z_j}}$$

Loss Function (Categorical Crossentropy):

$$\text{Loss} = - \sum y \log(\hat{y})$$

4. Algorithm Limitations

- Requires **large computational resources**
- Needs **large dataset** for better performance
- Training time is high
- Difficult to interpret (**black-box model**)
- Sensitive to hyperparameters
- May overfit if not properly regularized

5. Methodology / Workflow

Steps Involved:

1. **Dataset Loading**
Load MNIST dataset
2. **Data Preprocessing**
 - Normalize pixel values (0–1)
 - Reshape images for CNN
3. **Model Building**

- Add convolutional layers
- Add pooling layers
- Add fully connected layers
- 4. **Model Compilation**
 - Define optimizer (Adam)
 - Define loss function
- 5. **Model Training**
 - Train model using epochs
- 6. **Model Evaluation**
 - Calculate accuracy and loss
- 7. **Visualization**
 - Plot accuracy and loss graphs

6. Performance Analysis

The model's performance was evaluated using **Training vs. Validation** metrics over 5 epochs. The results are summarized below:

Key Metrics

- **Final Accuracy:** The model achieved a validation accuracy of approximately **98%**, indicating high-precision digit classification.
- **Final Loss:** The categorical crossentropy loss decreased significantly, stabilizing near **0.09**, which confirms efficient model convergence.

7. Hyperparameter Tuning

Hyperparameters Used:

- Epochs: 5
- Batch Size: 32
- Filters: 32, 64
- Kernel Size: (3,3)
- Optimizer: Adam

Impact:

- Increasing epochs improved accuracy
- Proper filter size improved feature extraction
- Adam optimizer ensured faster convergence

Code:

```
# Step 1: Import Libraries
import tensorflow as tf
from tensorflow.keras import layers, models
import matplotlib.pyplot as plt

# Step 2: Load Dataset (MNIST)
(X_train, y_train), (X_test, y_test) = tf.keras.datasets.mnist.load_data()

# Step 3: Preprocessing
plt.figure(figsize=(6,6))
for i in range(9):
    plt.subplot(3,3,i+1)
    plt.imshow(X_train[i], cmap='gray')
    plt.title(str(y_train[i]))
    plt.axis('off')
plt.show()

# Step 4: Preprocessing
X_train = X_train / 255.0
X_test = X_test / 255.0

# Reshape for CNN
X_train = X_train.reshape(-1,28,28,1)
X_test = X_test.reshape(-1,28,28,1)

# Step 5: Build CNN Model
model = models.Sequential()

model.add(layers.Conv2D(32, (3,3), activation='relu', input_shape=(28,28,1)))
model.add(layers.MaxPooling2D((2,2)))

model.add(layers.Conv2D(64, (3,3), activation='relu'))
model.add(layers.MaxPooling2D((2,2)))

model.add(layers.Flatten())
model.add(layers.Dense(64, activation='relu'))
model.add(layers.Dense(10, activation='softmax'))

# Step 6: Compile Model
model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])

# Step 7: Train Model
history = model.fit(
    X_train, y_train,
```

```
epochs=5,
validation_data=(X_test, y_test)
)

# Step 8: Evaluate Model
loss, accuracy = model.evaluate(X_test, y_test)
print("Accuracy:", accuracy)

# Step 9: Plot Accuracy Graph
plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])


plt.title('Model Accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend(['Training Accuracy', 'Validation Accuracy'])

plt.show()

# Step 10: Plot Loss Graph
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])

plt.title('Model Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend(['Training Loss', 'Validation Loss'])

plt.show()
```

Output:

Step 7: Train Model


```
history = model.fit(  
    X_train, y_train,  
    epochs=5,  
    validation_data=(X_test, y_test)  
)
```


```
Epoch 1/5  
1875/1875 ————— 65s 35ms/step - accuracy: 0.7726 - loss: 0.6931 - val_accuracy: 0.9055 - val_loss: 0.3072  
Epoch 2/5  
1875/1875 ————— 81s 34ms/step - accuracy: 0.9245 - loss: 0.2528 - val_accuracy: 0.9485 - val_loss: 0.1761  
Epoch 3/5  
1875/1875 ————— 62s 33ms/step - accuracy: 0.9532 - loss: 0.1550 - val_accuracy: 0.9654 - val_loss: 0.1120  
Epoch 4/5  
1875/1875 ————— 66s 35ms/step - accuracy: 0.9652 - loss: 0.1120 - val_accuracy: 0.9695 - val_loss: 0.0976  
Epoch 5/5  
1875/1875 ————— 82s 35ms/step - accuracy: 0.9717 - loss: 0.0902 - val_accuracy: 0.9770 - val_loss: 0.0753
```

Step 8: Evaluate Model

```
loss, accuracy = model.evaluate(X_test, y_test)  
print("Accuracy:", accuracy)
```

```
313/313 ————— 3s 10ms/step - accuracy: 0.9770 - loss: 0.0753  
Accuracy: 0.9769999980926514
```





Performance Analysis & Graph Interpretation

The training process was conducted over **5 epochs** using the **Adam optimizer**. The resulting plots illustrate the model's learning efficiency and generalization capability.

A. Accuracy Analysis (Model Accuracy)

- **Trend:** Both training and validation accuracy show a consistent **upward trend**.
- **Final Accuracy:** The model achieved a training accuracy of approximately **97.2%** and a validation accuracy of **97.8%**.
- **Generalization:** The validation accuracy remains slightly higher than the training accuracy throughout, indicating that the model is **generalizing exceptionally well** to unseen data without signs of overfitting.

B. Loss Analysis (Model Loss)

- **Trend:** The loss for both sets shows a **steep exponential decay**, which is the ideal behavior for a converging neural network.

- **Convergence:** The training loss dropped from **0.70** to below **0.10**.
- **Optimization:** The smooth, downward slope of the validation loss confirms that the learning rate and batch size (32) were appropriately tuned, as there are no large oscillations or "spikes" in the data.

Conclusion:

The design and implementation of a **Convolutional Neural Network (CNN)** for the **MNIST Handwritten Digit Dataset** were successfully executed. By utilizing a layered architecture of **Conv2D** for feature extraction and **MaxPooling** for spatial reduction, the model achieved a high-performance classification rate of nearly **98%**.

The experiment demonstrates that:

1. **CNN Superiority:** Convolutional layers are far more effective than standard dense layers for image data as they preserve the spatial relationship between pixels.
2. **Effective Convergence:** Through proper normalization (scaling pixels to 0–1) and the use of the **ReLU** activation function, the model reached high accuracy in a very short training window (5 epochs).
3. **Model Robustness:** The minimal gap between the training and validation curves confirms a robust fit, proving the model is ready for real-world digit recognition tasks.